

```
In [2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [3]: #import interactive library
from ipywidgets import interact
```

```
In [4]: #balancing the target variable
from imblearn.over_sampling import RandomOverSampler
```

```
In [5]: #variable scaling libraries
from sklearn.preprocessing import StandardScaler
```

```
In [6]: #skLearn_libraries
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
```

```
In [7]: #feature_selections
from sklearn.feature_selection import f_classif
from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import chi2
```

```
In [28]: #evaluation_metrics
from sklearn import metrics
from sklearn.metrics import f1_score
from sklearn.metrics import accuracy_score
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix

#import hyperparameter tuning library
from sklearn.model_selection import GridSearchCV

#other libraries
import math
from collections import Counter
```

```
In [9]: data=pd.read_csv(r"C:\Users\SENAPATHI REDDY.K\Downloads\waterQuality1.csv")
```

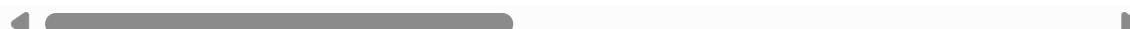
EDA_ANALYSIS

```
In [10]: data.head()
```

Out[10]:

	aluminium	ammonia	arsenic	barium	cadmium	chloramine	chromium	copper	f
0	1.65	9.08	0.04	2.85	0.007	0.35	0.83	0.17	
1	2.32	21.16	0.01	3.31	0.002	5.28	0.68	0.66	
2	1.01	14.02	0.04	0.58	0.008	4.24	0.53	0.02	
3	1.36	11.33	0.04	2.96	0.001	7.23	0.03	1.66	
4	0.92	24.33	0.03	0.20	0.006	2.67	0.69	0.57	

5 rows × 21 columns



In [11]: `data.shape`

Out[11]: (7999, 21)

In [12]: `data.columns`

Out[12]: Index(['aluminium', 'ammonia', 'arsenic', 'barium', 'cadmium', 'chloramine', 'chromium', 'copper', 'flouride', 'bacteria', 'viruses', 'lead', 'nitrates', 'nitrites', 'mercury', 'perchlorate', 'radium', 'selenium', 'silver', 'uranium', 'is_safe'],
dtype='object')

```
In [13]: data.isnull().sum()
```

```
Out[13]: aluminium      0
ammonia          0
arsenic          0
barium           0
cadmium          0
chloramine       0
chromium         0
copper            0
flouride         0
bacteria         0
viruses          0
lead              0
nitrates          0
nitrites          0
mercury           0
perchlorate       0
radium            0
selenium          0
silver             0
uranium           0
is_safe            0
dtype: int64
```

```
In [14]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7999 entries, 0 to 7998
Data columns (total 21 columns):
 #   Column      Non-Null Count  Dtype  
 ---  --          -----          float64
 0   aluminium   7999 non-null   float64
 1   ammonia     7999 non-null   object 
 2   arsenic     7999 non-null   float64
 3   barium      7999 non-null   float64
 4   cadmium     7999 non-null   float64
 5   chloramine  7999 non-null   float64
 6   chromium    7999 non-null   float64
 7   copper      7999 non-null   float64
 8   fluoride    7999 non-null   float64
 9   bacteria    7999 non-null   float64
 10  viruses     7999 non-null   float64
 11  lead        7999 non-null   float64
 12  nitrates    7999 non-null   float64
 13  nitrites    7999 non-null   float64
 14  mercury     7999 non-null   float64
 15  perchlorate 7999 non-null   float64
 16  radium      7999 non-null   float64
 17  selenium    7999 non-null   float64
 18  silver      7999 non-null   float64
 19  uranium     7999 non-null   float64
 20  is_safe     7999 non-null   object 
```

```
dtypes: float64(19), object(2)
memory usage: 1.3+ MB
```

```
In [15]: #drop missing values
data['ammonia'].value_counts()['#NUM!']
data = data[data['ammonia'].str.contains('#NUM!') == False]
```

```
In [16]: #convert both columns to numeric data type
data['ammonia'] = pd.to_numeric(data['ammonia'])
data['is_safe'] = pd.to_numeric(data['is_safe'])
```

```
In [17]: #statistic metrics for continuous variables without scientific notation
data.describe().apply(lambda s: s.apply(lambda x: format(x, 'f')))
```

Out[17]:

	aluminium	ammonia	arsenic	barium	cadmium	chloramine	
count	7996.000000	7996.000000	7996.000000	7996.000000	7996.000000	7996.000000	7
mean	0.666396	14.278212	0.161477	1.567928	0.042803	2.177589	
std	1.265323	8.878930	0.252632	1.216227	0.036049	2.567210	
min	0.000000	-0.080000	0.000000	0.000000	0.000000	0.000000	
25%	0.040000	6.577500	0.030000	0.560000	0.008000	0.100000	
50%	0.070000	14.130000	0.050000	1.190000	0.040000	0.530000	
75%	0.280000	22.132500	0.100000	2.482500	0.070000	4.240000	
max	5.050000	29.840000	1.050000	4.940000	0.130000	8.680000	

8 rows × 21 columns



In [18]:

```
#create a list of numerical features and plot them
list_of_num_features = data.loc[:, data.columns != 'is_safe']
palette_features = ['#E68753', '#409996']
sns.set(rc={'axes.facecolor':'#ECECEC'}) #background color of all plots

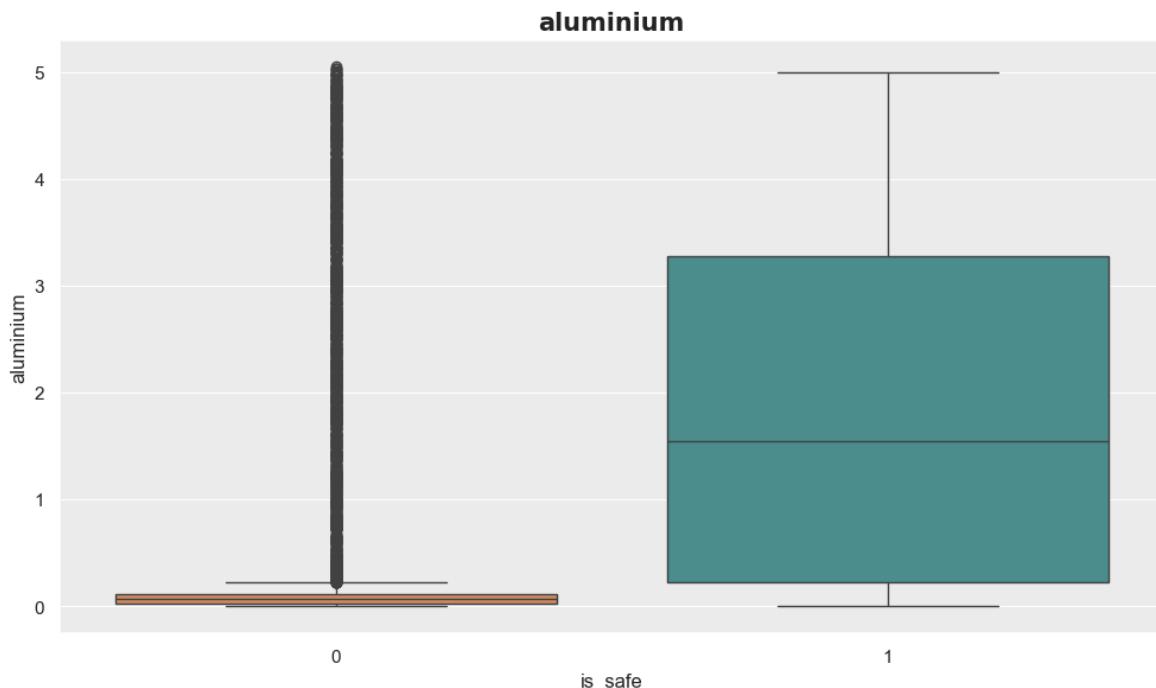
for feature in list_of_num_features:
```

```
plt.figure(figsize=(12,6.5))
plt.title(feature, fontsize=15, fontweight='bold', fontname='Helvetica', ha=
ax = sns.boxplot(
x='is_safe',
y=list_of_num_features[feature],
data=data,
hue='is_safe',
palette=palette_features,
legend=False
)
```

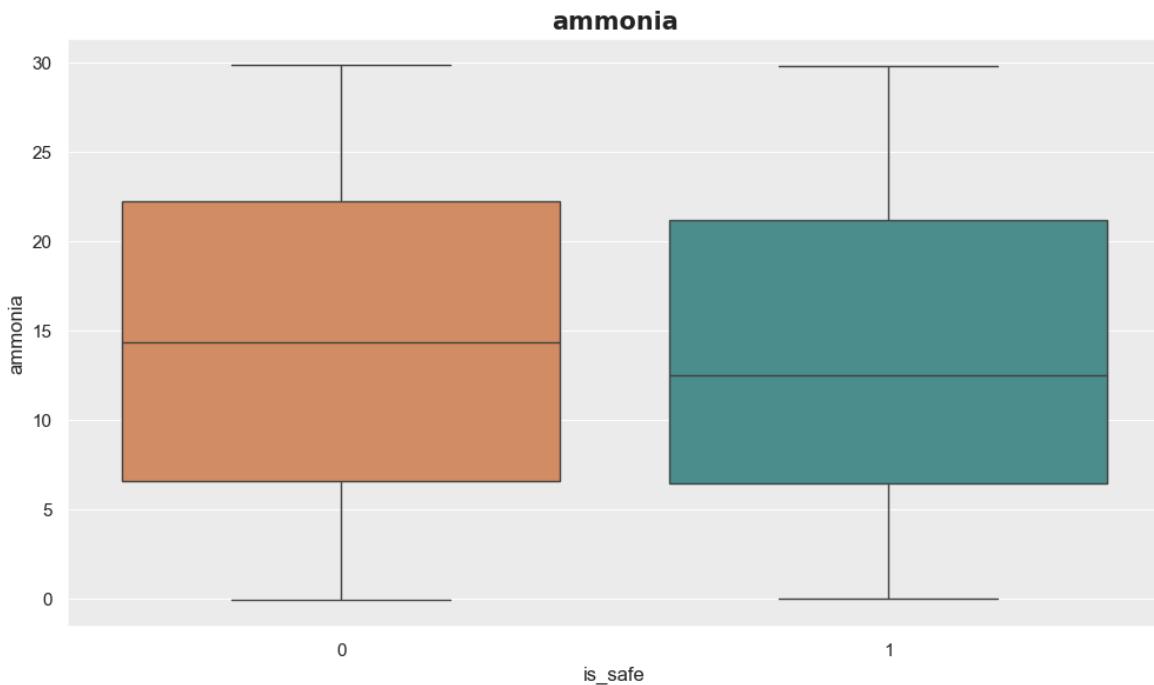
```
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.
```

```
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.
```

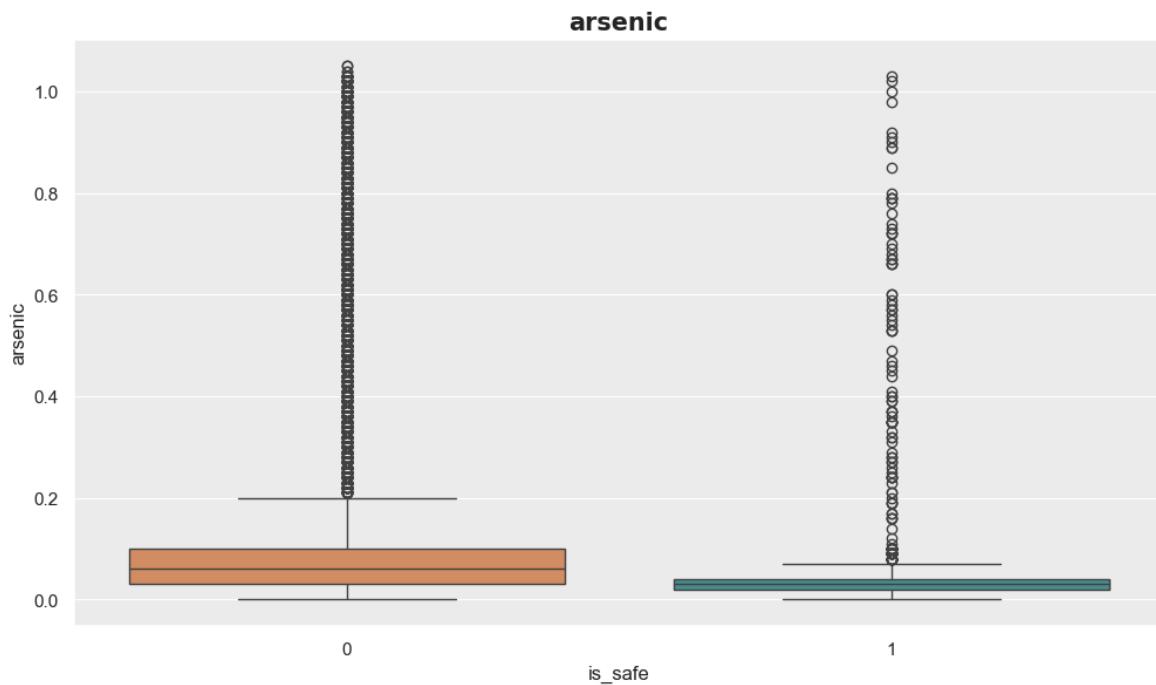
```
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.
```



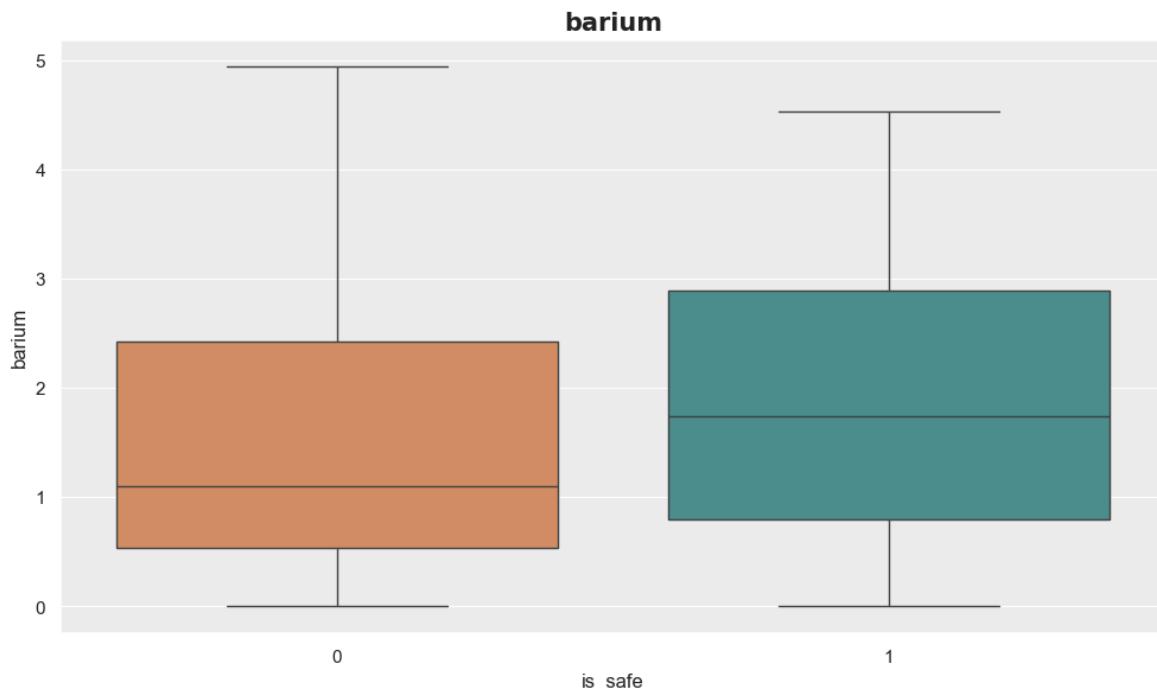
```
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.
```



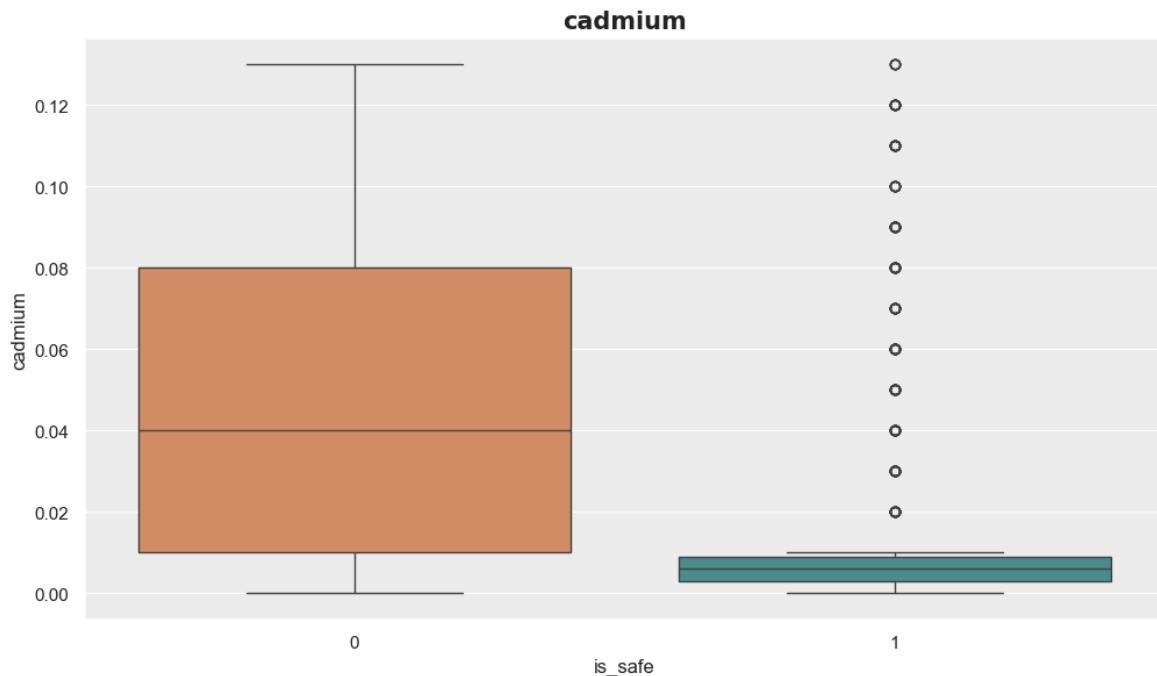
```
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.
```



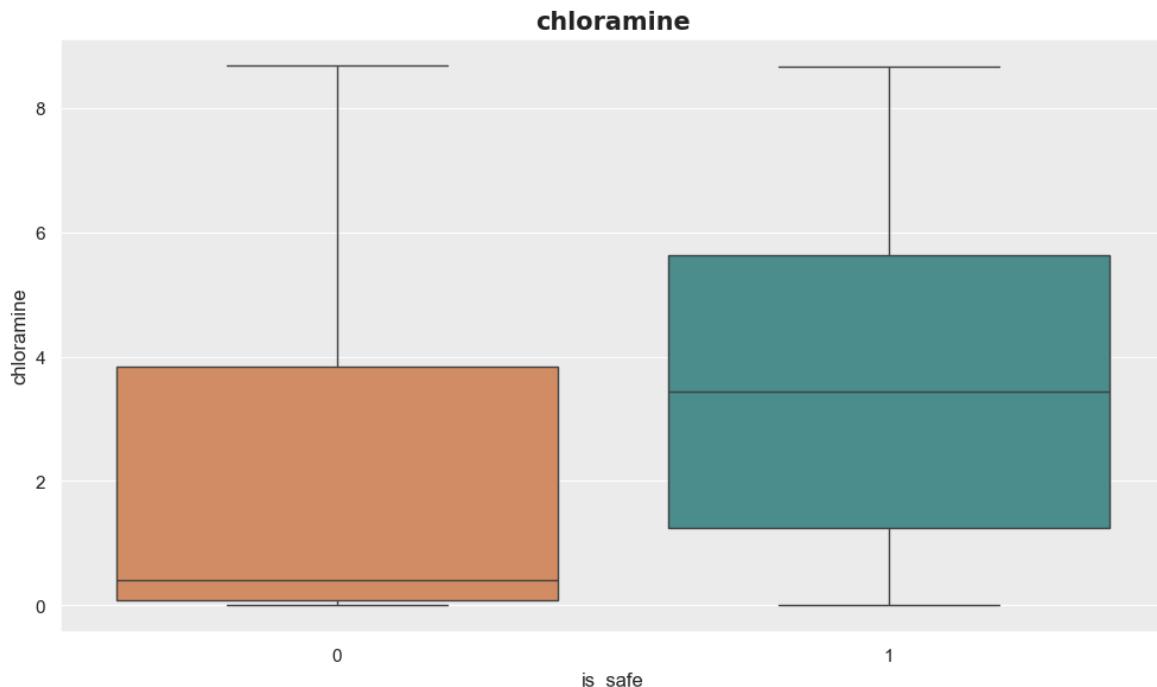
```
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.
```



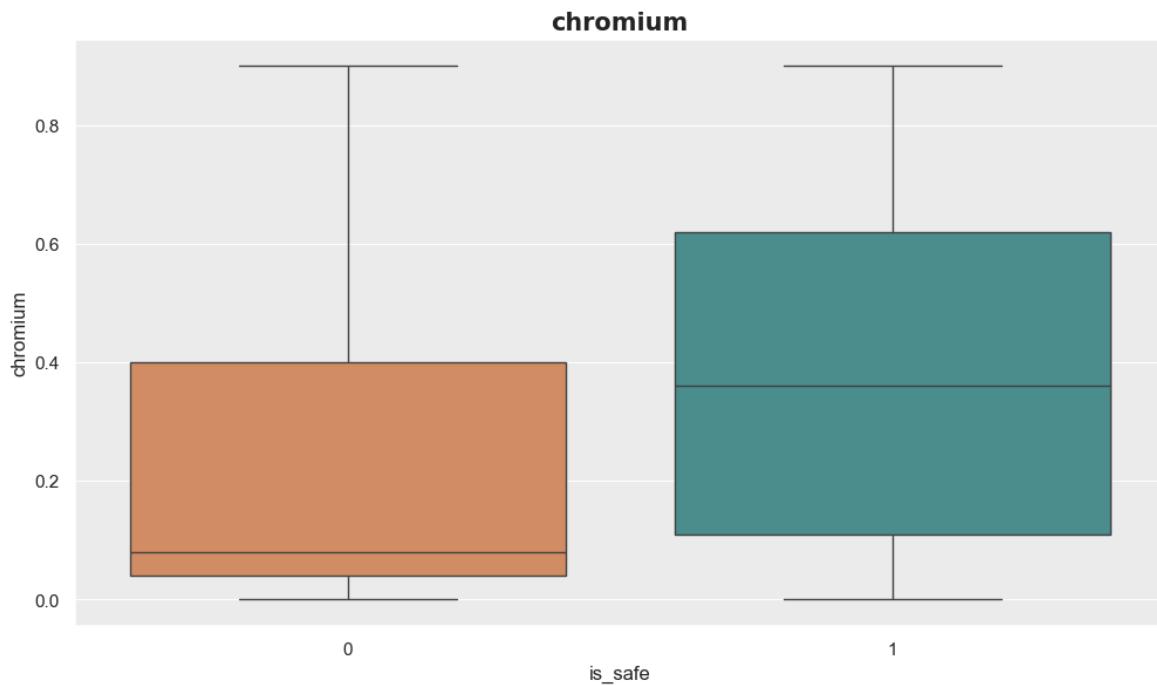
```
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.
```



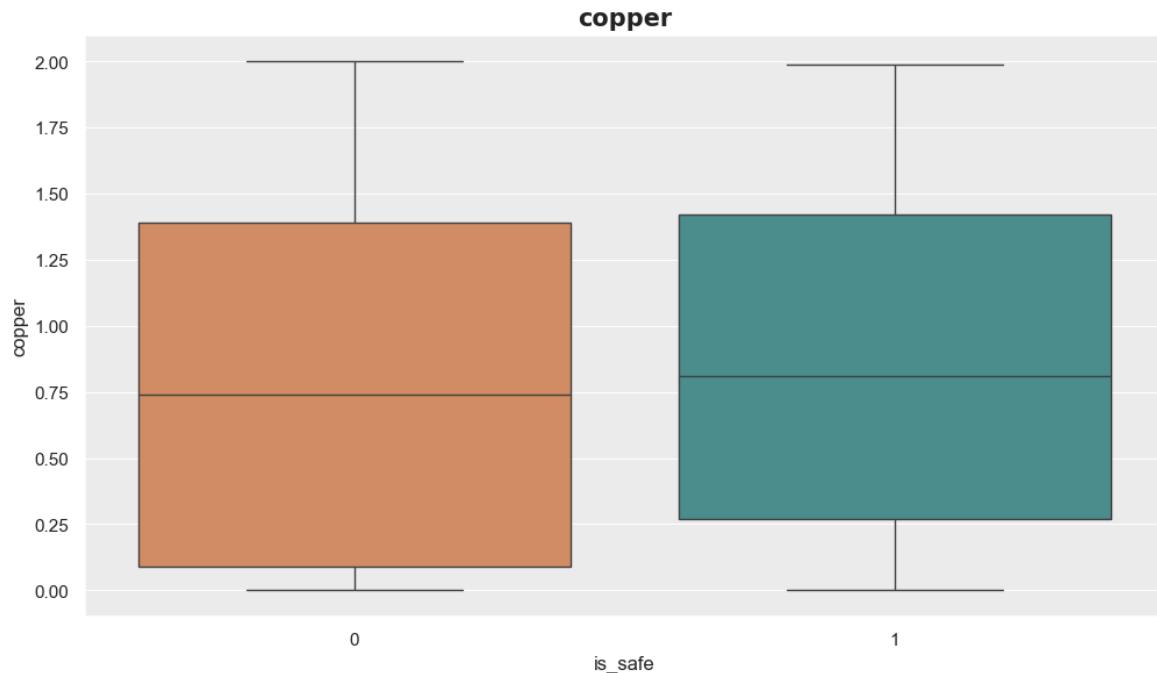
```
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.
```



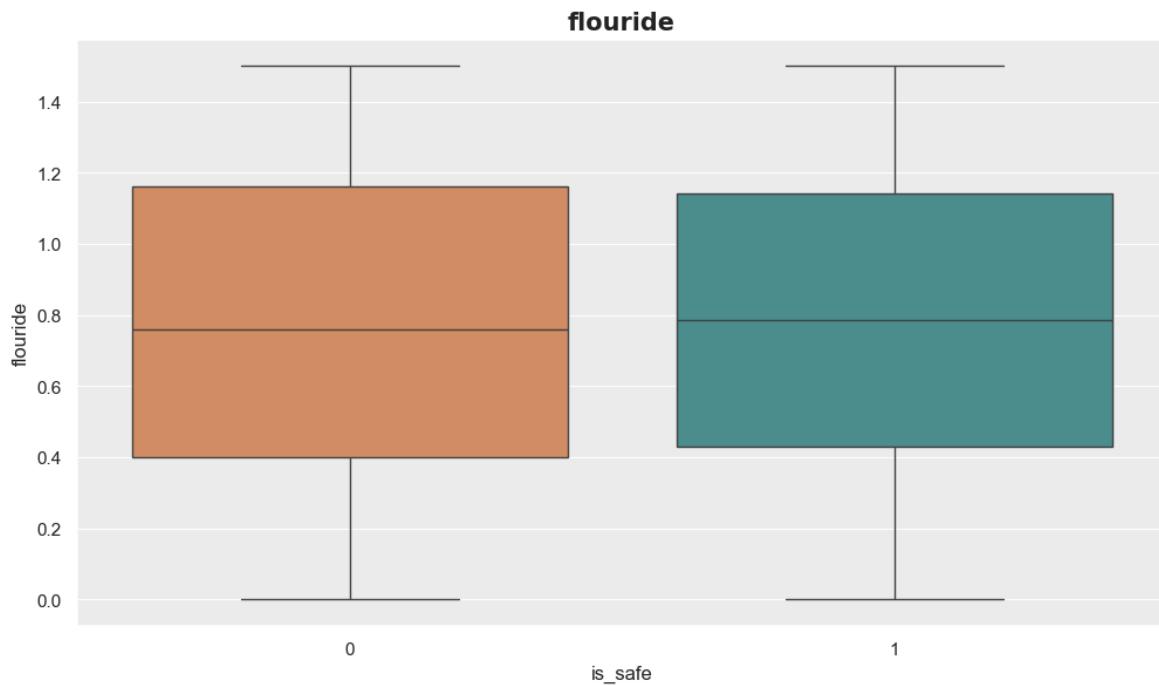
```
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.
```



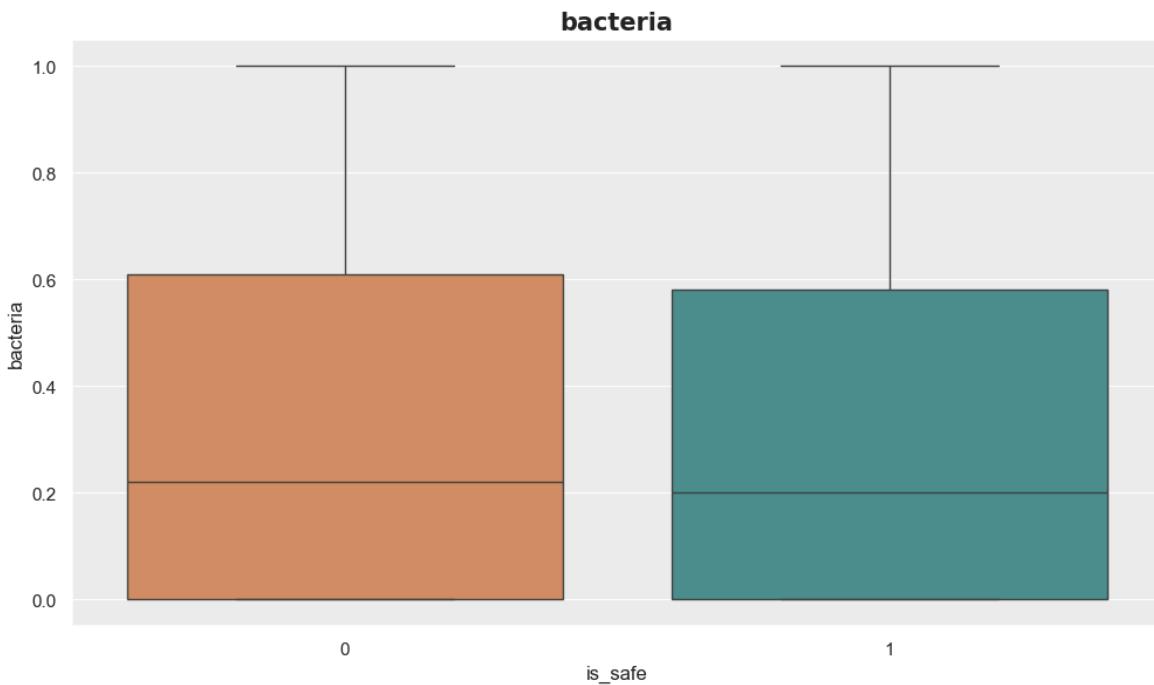
```
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.
```



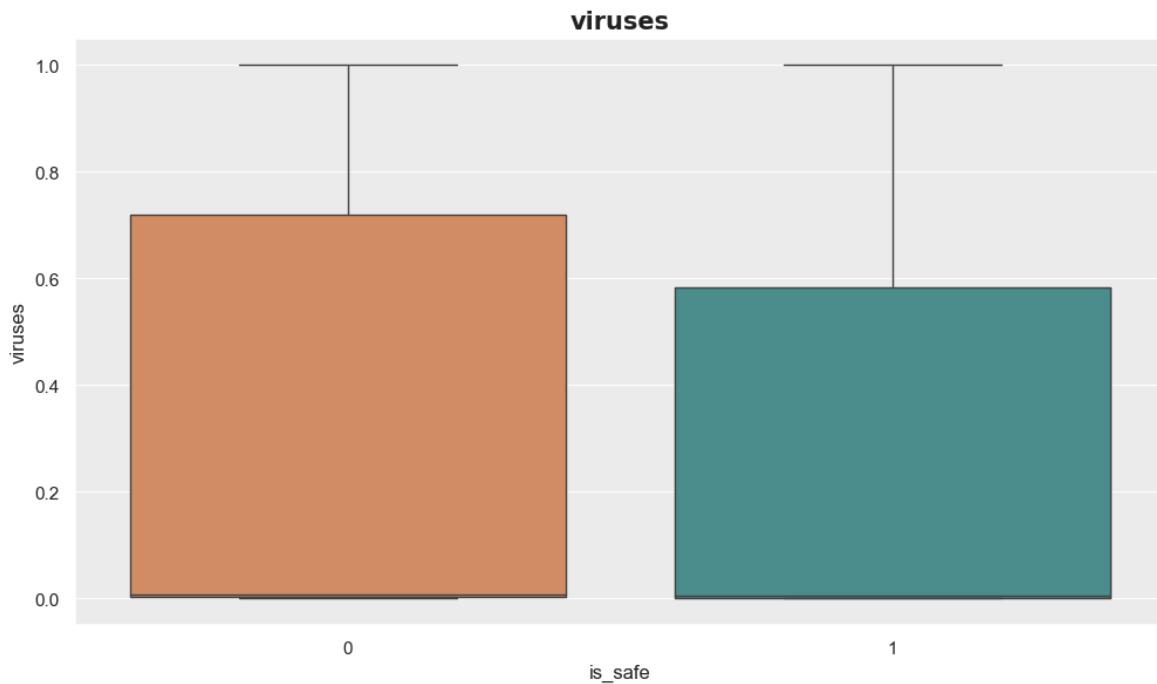
```
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.
```



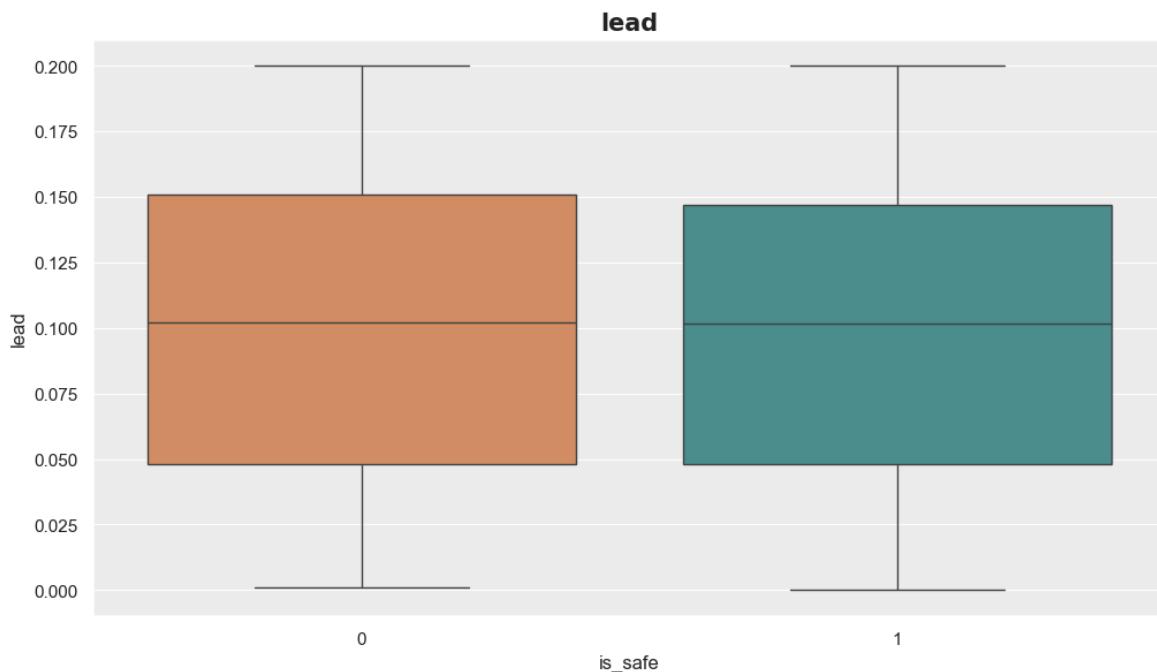
```
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.
```



```
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.
```



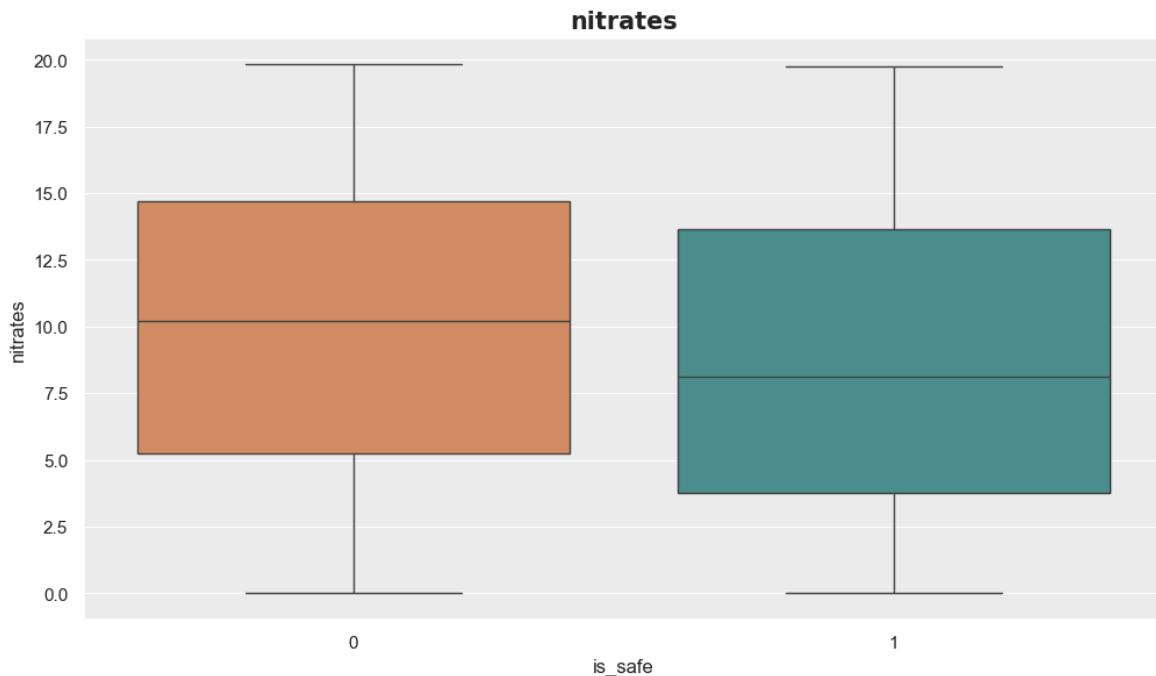
```
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.
```



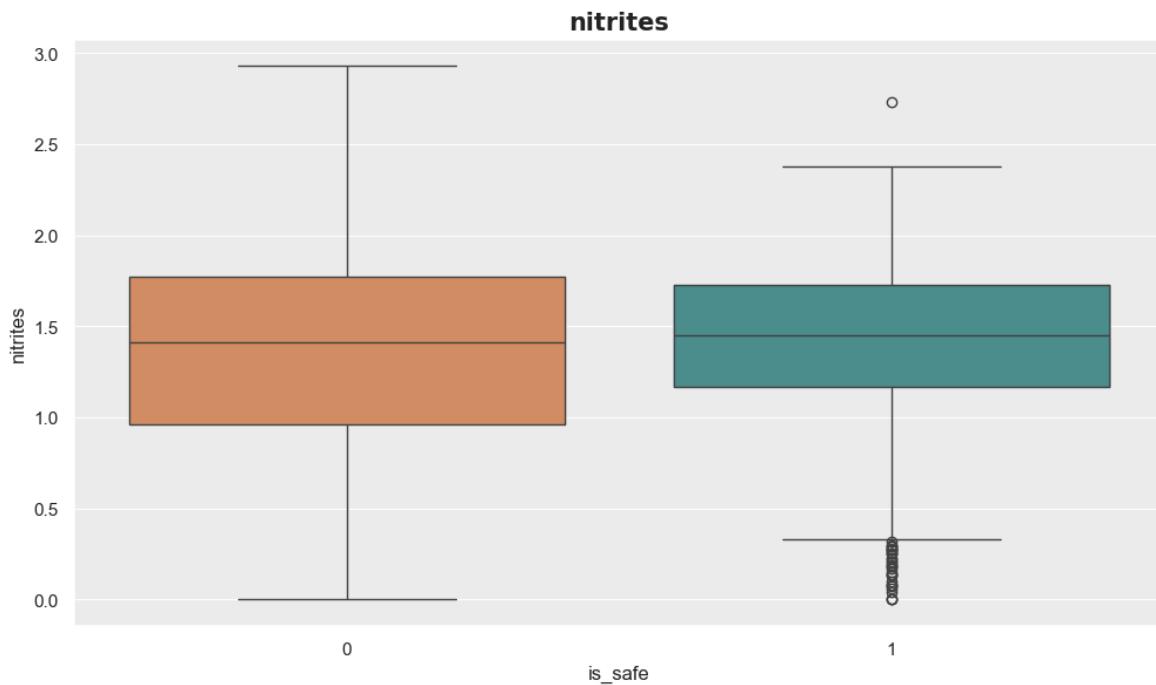
```
findfont: Font family 'Helvetica' not found.
```

```
findfont: Font family 'Helvetica' not found.
```

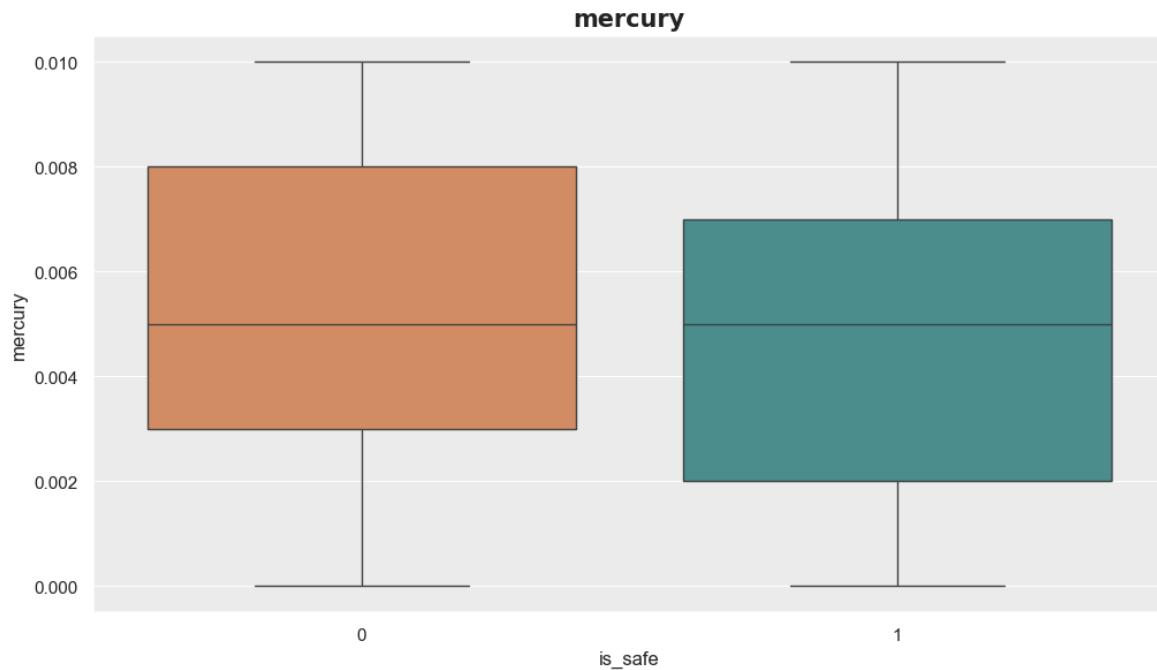
```
findfont: Font family 'Helvetica' not found.
```



```
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.
```



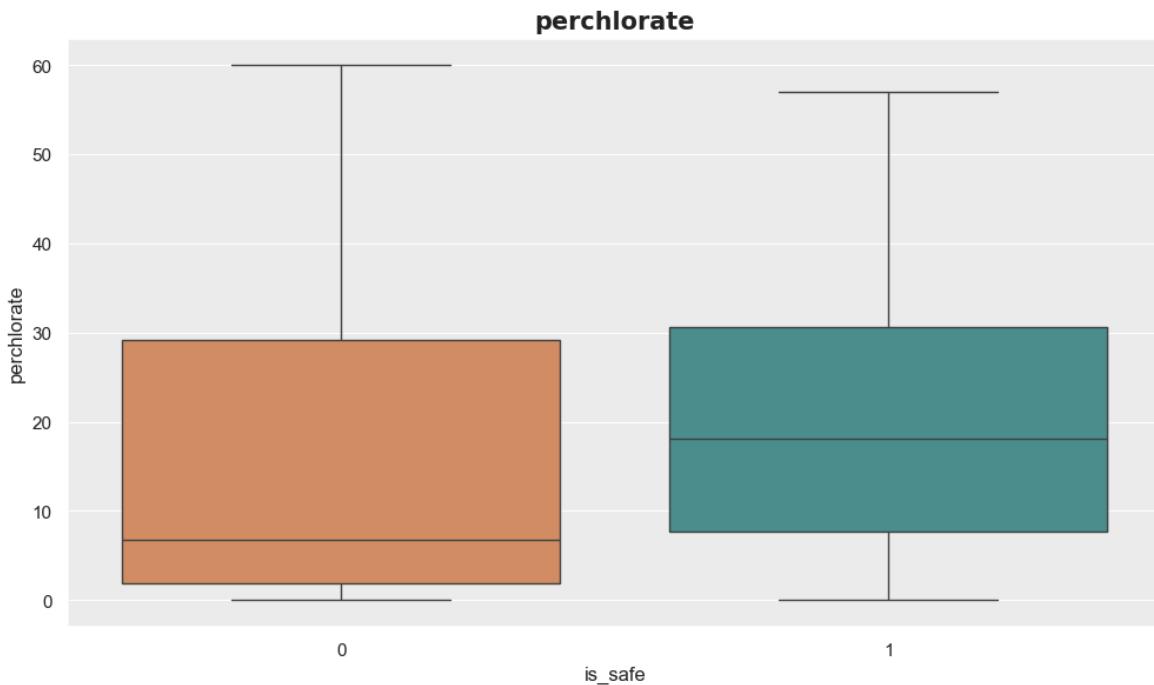
```
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.
```



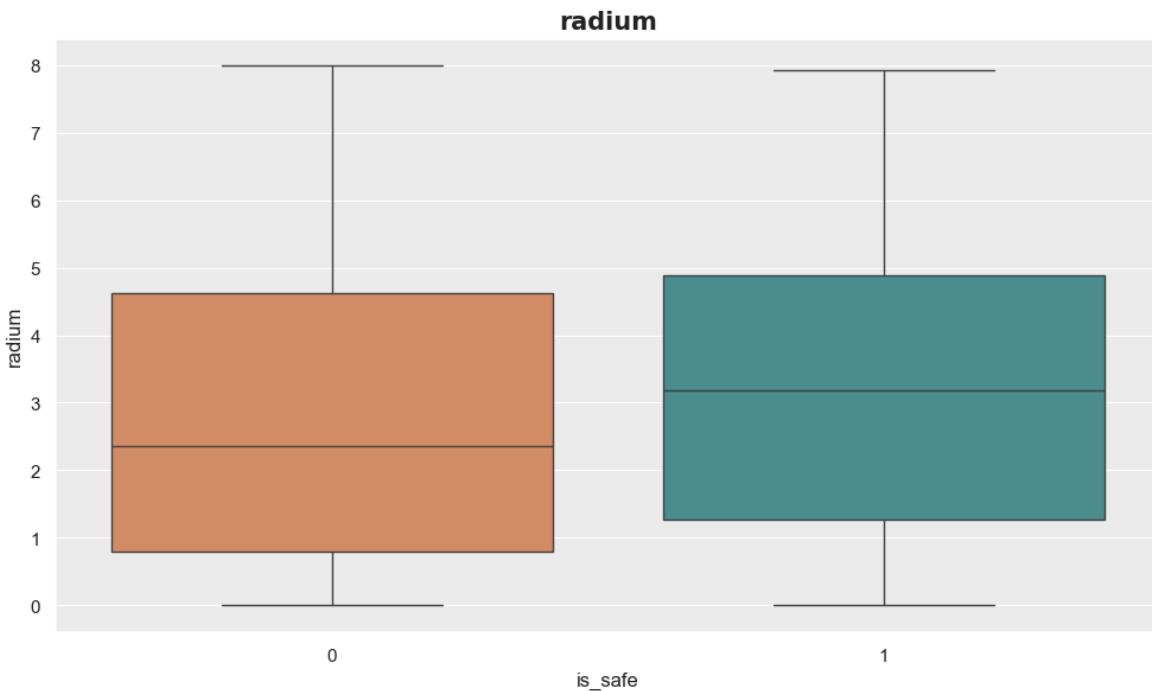
```
findfont: Font family 'Helvetica' not found.
```

```
findfont: Font family 'Helvetica' not found.
```

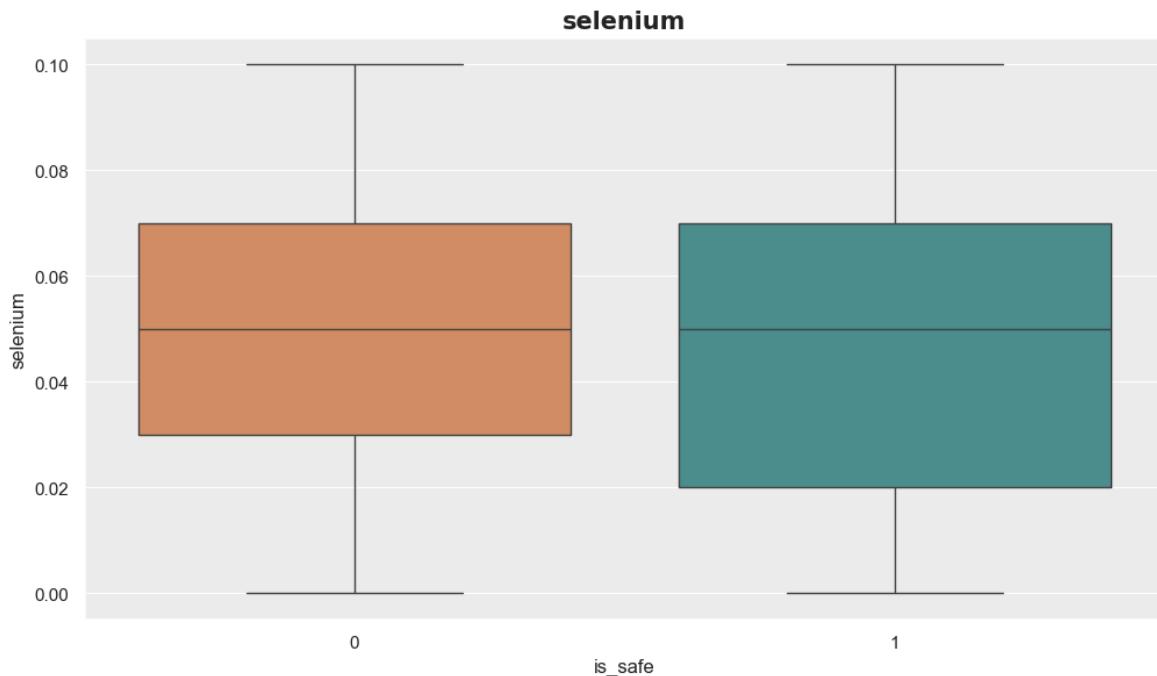
```
findfont: Font family 'Helvetica' not found.
```



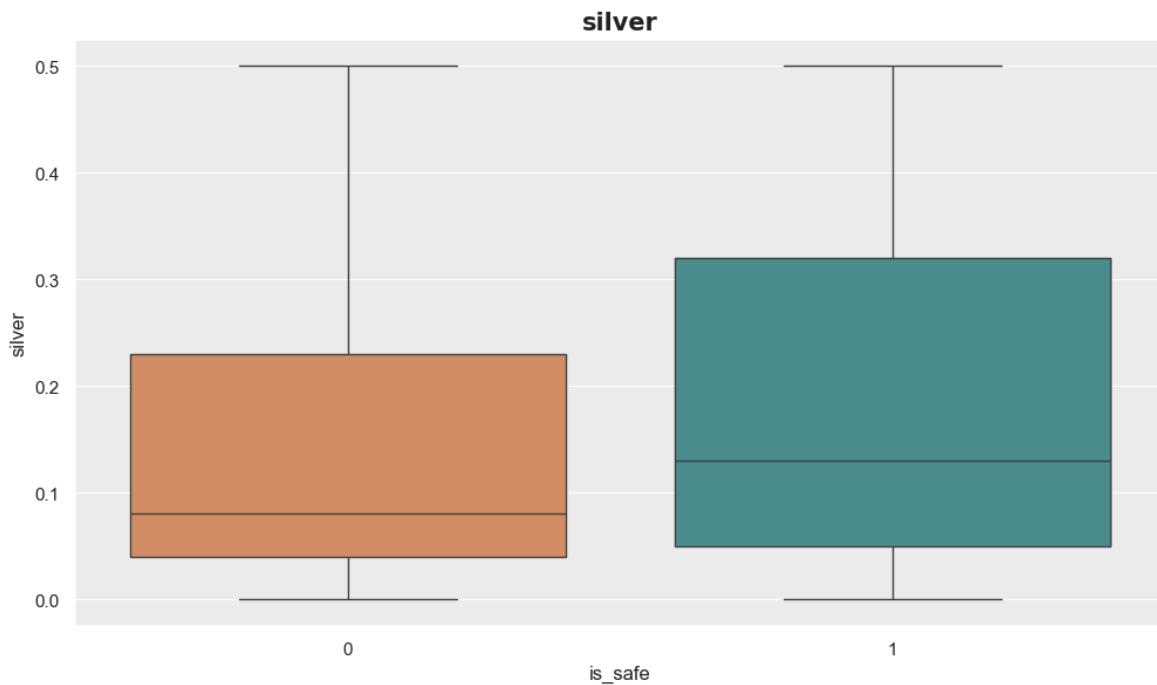
```
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.
```



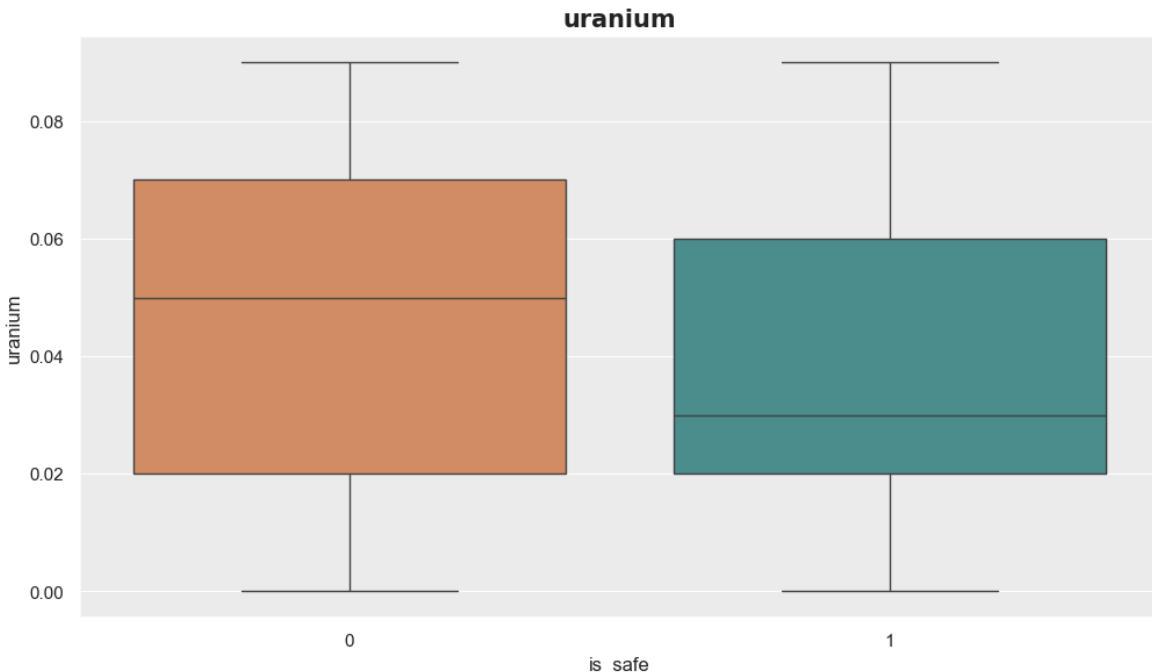
```
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.
```



```
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.
```



```
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.  
findfont: Font family 'Helvetica' not found.
```



```
In [20]: #balance the target variable
columns = data.columns
columns = [c for c in columns if c not in ['is_safe']]
y = data['is_safe'] #prior target variable
x = data[columns] #prior features
```

```
ros = RandomOverSampler(sampling_strategy='minority')
x, y = ros.fit_resample(x, y) #y_train as balanced target variable
print(f"Imbalanced target class: {(y)}\n\nBalanced target class: {Counter(y)}\n"
print(x.shape[0] - data.shape[0], 'new random picked points')
```

```
Imbalanced target class: 0           1
1      1
2      0
3      1
4      1
..
14163   1
14164   1
14165   1
14166   1
14167   1
Name: is_safe, Length: 14168, dtype: int64
```

```
Balanced target class: Counter({1: 7084, 0: 7084})
```

```
6172 new random picked points
```

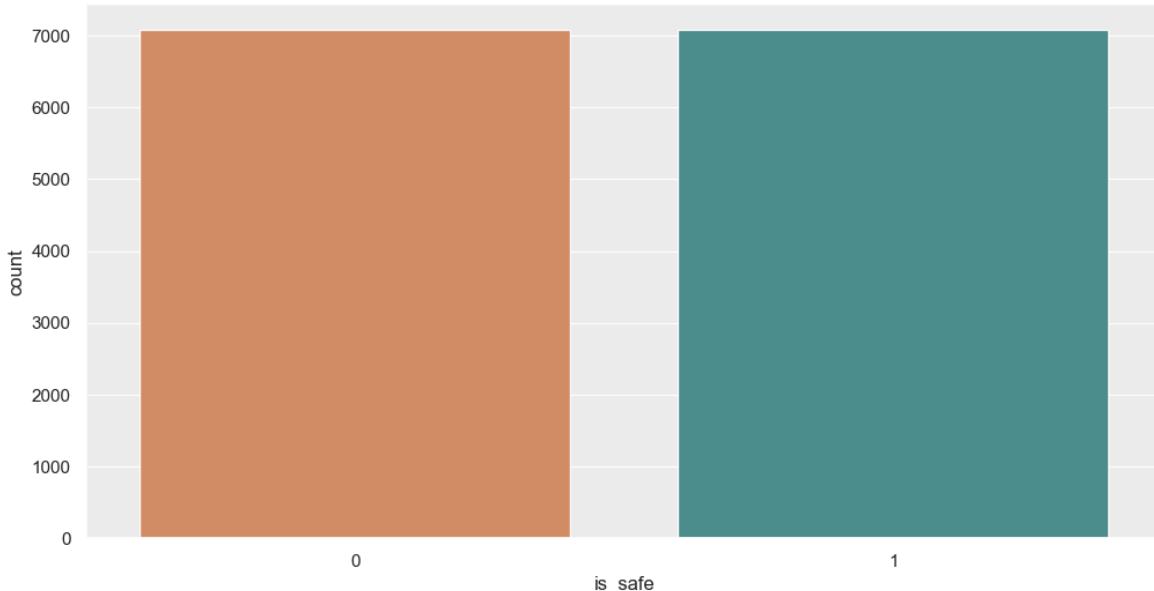
```
In [21]: #plot the balanced target variable
sns.set(rc={'axes.facecolor':'#ECECEC'}) #background color of plot
plt.figure(figsize=(12,6))
plt.title("Balanced target variable", fontsize=15, fontweight='bold', fontname='
```

```
ax = sns.countplot(x=y, data=data, palette=palette_features)
plt.show()
```

C:\Users\SENAPATHI REDDY.K\AppData\Local\Temp\ipykernel_18828\1791969356.py:5: FutureWarning:

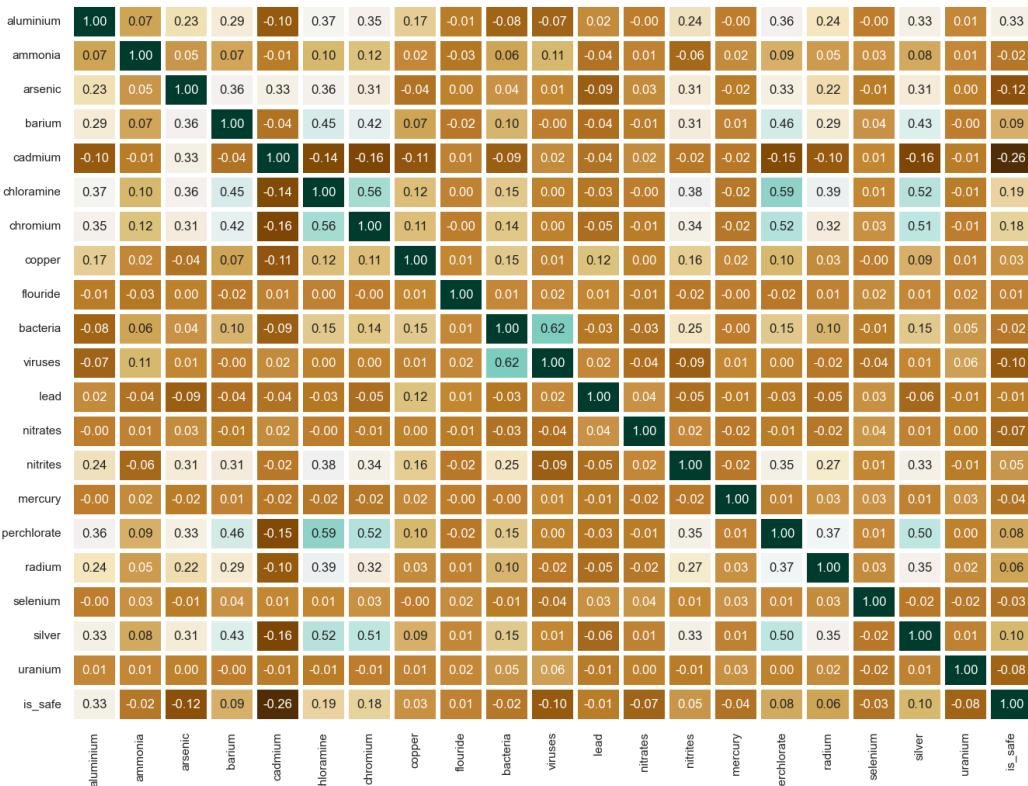
Passing `palette` without assigning `hue` is deprecated and will be removed in v 0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
ax = sns.countplot(x=y, data=data, palette=palette_features)
findfont: Font family 'Helvetica' not found.
```

Balanced target variable

```
In [22]: # plotting correlation matrix to notice relationships or lack of it between vari  
corr = data.corr()  
  
plt.figure(figsize = (20, 12))  
sns.heatmap(corr, xticklabels = corr.columns, yticklabels = corr.columns, linewidths=1)  
plt.show()
```

KNN_Regressor



```
In [23]: #preparation of training and testing data  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=42)
```

```
In [24]: x_train.shape,x_test.shape,y_train.shape,y_test.shape
```

```
Out[24]: ((11334, 20), (2834, 20), (11334,), (2834,))
```

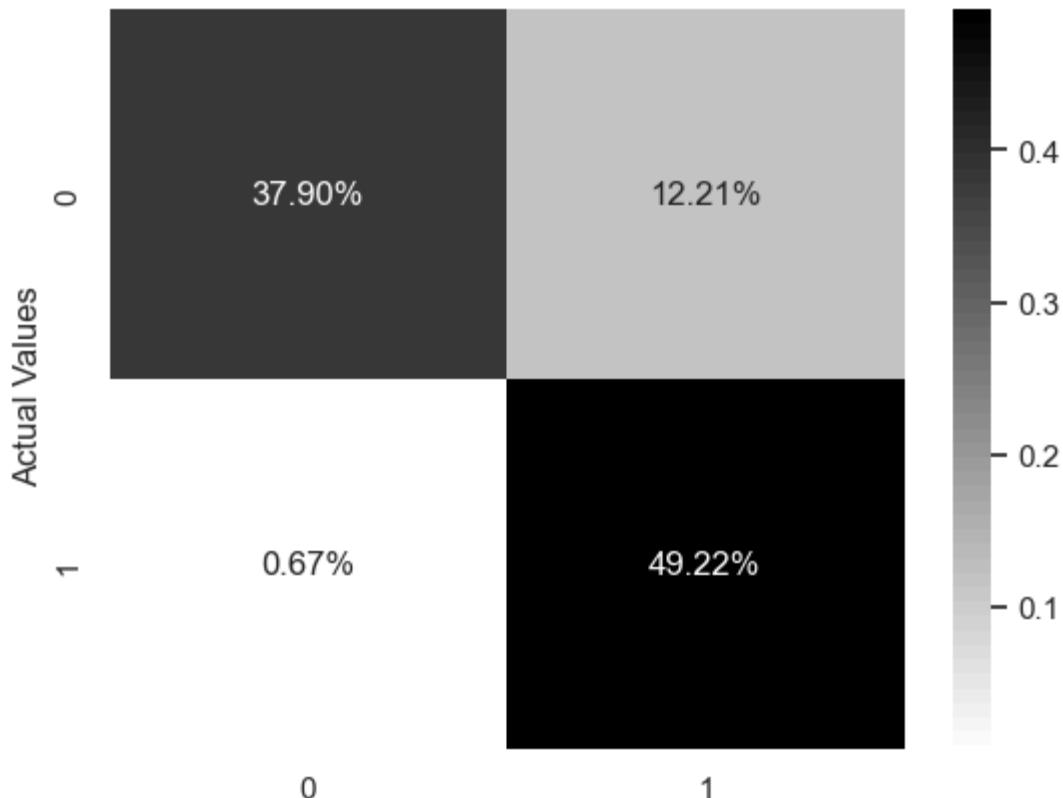
```
In [26]: knnc=KNeighborsClassifier()  
knnc.fit(x_train,y_train)  
y_pred=knnc.predict(x_test)
```

```
In [29]: cf_matrix=confusion_matrix(y_test,y_pred)  
print(cf_matrix)
```

```
[[1074  346]  
 [ 19 1395]]
```

```
In [31]: ax = sns.heatmap(cf_matrix/np.sum(cf_matrix), annot=True, fmt=' .2%', cmap='binary')  
ax.set_title('K-NN Confusion Matrix\n\n');  
ax.set_xlabel('\nPredicted Values')  
ax.set_ylabel('Actual Values ');  
ax.xaxis.set_ticklabels(['0','1'])  
ax.yaxis.set_ticklabels(['0','1'])  
plt.show()
```

K-NN Confusion Matrix



Predicted Values

In [32]:

```
print(classification_report(y_test,y_pred))
print("accuracy_score",accuracy_score(y_test,y_pred))
print("precision_score",precision_score(y_test,y_pred))
print("f1-score",f1_score(y_test,y_pred))
print("recall_score",recall_score(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.98	0.76	0.85	1420
1	0.80	0.99	0.88	1414
accuracy			0.87	2834
macro avg	0.89	0.87	0.87	2834
weighted avg	0.89	0.87	0.87	2834


```
accuracy_score 0.8712067748764997
precision_score 0.8012636415852958
f1-score 0.884310618066561
recall_score 0.9865629420084866
```

In [33]:

```
x_train,x_test,y_train,y_test = train_test_split(x, y,test_size=0.3, random_state=42)
x_train.shape,x_test.shape,y_train.shape,y_test.shape
```

Out[33]: ((9917, 20), (4251, 20), (9917,), (4251,))

```
In [38]: sc = StandardScaler()
sc.fit(x_train)
x_train_std = sc.transform(x_train)
x_test_std = sc.transform(x_test)
print('After standardizing our features, the first 5 rows of our data now look like')
print(pd.DataFrame(x_train_std).head())
```

After standardizing our features, the first 5 rows of our data now look like this:

	0	1	2	3	4	5	6	\
0	-0.747105	-1.287465	-0.515795	-1.229938	-0.879215	0.502308	0.900971	
1	0.470680	-1.600619	-0.426214	1.078939	-0.681569	1.303799	0.226259	
2	-0.148086	0.112650	-0.471004	-0.722313	-0.794510	-0.452578	0.510348	
3	1.168438	0.131938	-0.515795	0.563126	-0.681569	-0.425734	-0.199875	
4	-0.286321	-0.534081	-0.515795	1.463752	-0.879215	0.180178	-0.839076	

	7	8	9	10	11	12	13	\
0	-0.129825	1.104969	1.711217	1.587798	-0.037843	-1.600573	-0.134877	
1	0.429102	-0.396231	1.281916	1.207030	1.370857	0.496324	0.956896	
2	0.972504	-0.765758	-0.956581	-0.778405	1.701293	1.043028	1.033512	
3	0.444628	0.850920	-0.956581	-0.778405	0.379550	-1.233706	0.037508	
4	0.040958	1.174255	0.239328	-0.764807	-1.446543	0.361447	0.056662	

	14	15	16	17	18	19	
0	0.330781	0.496439	-0.905075	0.728547	-0.025737	1.425689	
1	0.994352	1.101963	0.319034	0.382140	1.116711	1.052404	
2	-1.328144	-0.611256	0.766137	0.035734	-0.832171	0.679118	
3	-0.664574	-0.173528	1.933815	1.767766	-1.033780	-0.440738	
4	0.994352	0.743269	1.530119	-1.349892	1.788740	-0.067453	

```
In [39]: # find best scored 10 features
select_feature = SelectKBest(f_classif, k=10).fit(x_train, y_train)
```

```
print('Score list:', select_feature.scores_)
print('Feature list:', x_train.columns)
```

```
Score list: [2.20931713e+03 6.18990165e+00 4.91501221e+02 1.81201115e+02
2.06393780e+03 9.72225517e+02 7.59077165e+02 1.49932412e+01
1.26787348e+00 1.17728614e+01 2.47876336e+02 3.85561481e+00
1.11624492e+02 5.19701896e+01 4.76808986e+01 1.56665064e+02
9.93507786e+01 1.91311693e+01 2.19988180e+02 1.57347696e+02]
Feature list: Index(['aluminium', 'ammonia', 'arsenic', 'barium', 'cadmium', 'chl
oramine',
       'chromium', 'copper', 'flouride', 'bacteria', 'viruses', 'lead',
       'nitrates', 'nitrites', 'mercury', 'perchlorate', 'radium', 'selenium',
       'silver', 'uranium'],
      dtype='object')
```

In [40]:

```
x_train_2 = select_feature.transform(x_train)
x_test_2 = select_feature.transform(x_test)

knnc = KNeighborsClassifier()
knnc = knnc.fit(x_train_2,y_train)

y_pred_knnc = knnc.predict(x_test_2)

cf_matrix_knnc = confusion_matrix(y_test, y_pred_knnc)
print(cf_matrix_knnc)

ax = sns.heatmap(cf_matrix_knnc/np.sum(cf_matrix_knnc), annot=True, fmt='.2%', c
```

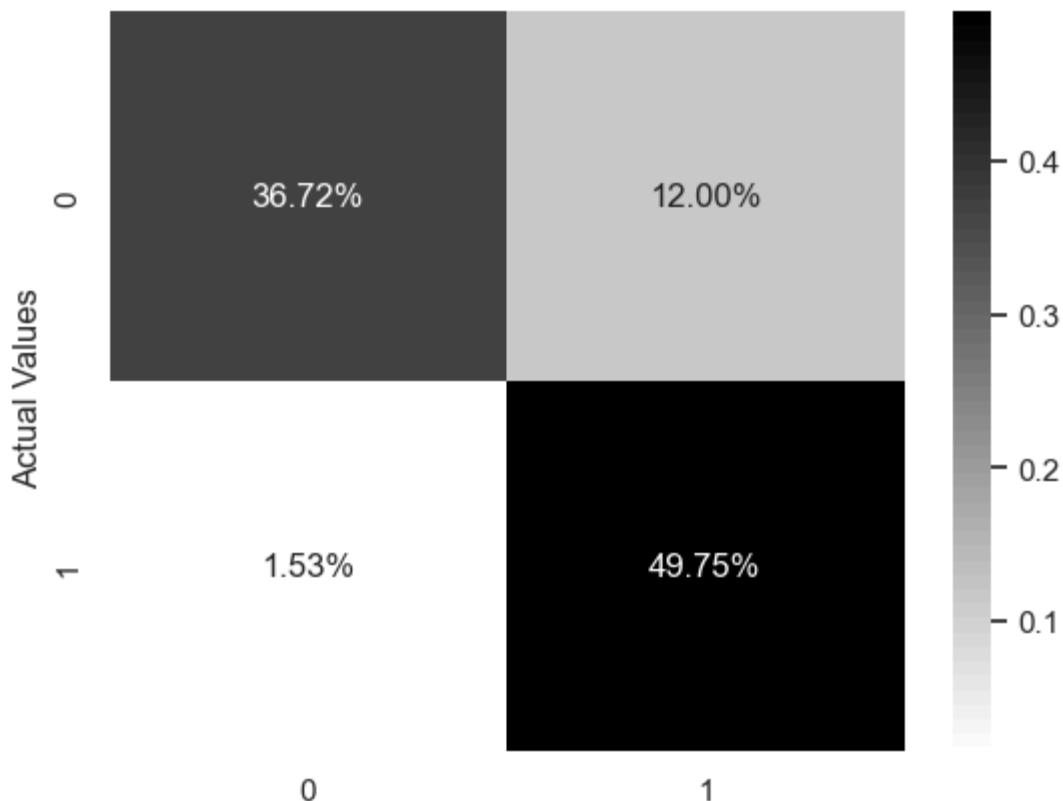
```
ax.set_title('K-NN Confusion Matrix\n\n');
ax.set_xlabel('\nPredicted Values')
ax.set_ylabel('Actual Values ');

## Ticket Labels - List must be in alphabetical order
ax.xaxis.set_ticklabels(['0','1'])
ax.yaxis.set_ticklabels(['0','1'])

plt.show()
```

```
[[1561  510]
 [ 65 2115]]
```

K-NN Confusion Matrix



Predicted Values

```
In [41]: #evaluation of K-NN: metrics pivot chart
print(classification_report(y_test, y_pred_knnc))

print('Accuracy Score : ' + str(round(accuracy_score(y_test,y_pred_knnc),3)))
print('Precision Score : ' + str(round(precision_score(y_test,y_pred_knnc, pos_label=0),3)))
print('Recall Score : ' + str(round(recall_score(y_test,y_pred_knnc, pos_label=0),3)))
print('F-Score : ' + str(round(f1_score(y_test,y_pred_knnc, pos_label=0),3)))
```

	precision	recall	f1-score	support
0	0.96	0.75	0.84	2071
1	0.81	0.97	0.88	2180
accuracy			0.86	4251
macro avg	0.88	0.86	0.86	4251
weighted avg	0.88	0.86	0.86	4251

Accuracy Score : 0.865
Precision Score : 0.96
Recall Score : 0.754
F-Score : 0.844

```
In [42]: param_grid = {'n_neighbors':np.arange(1,40), 'metric':['euclidean', 'manhattan'],
knn = KNeighborsClassifier()
knn_cv = GridSearchCV(knn,param_grid,cv=5)
knn_cv.fit(x_train_2,y_train)

y_pred_knn_cv = knn_cv.predict(X_test_2)

print("Best Score: " + str(knn_cv.best_score_))
print("Best Parameters: " + str(knn_cv.best_params_))

cf_matrix_knn_cv = confusion_matrix(y_test, y_pred_knn_cv)
print(cf_matrix_knn_cv)

ax = sns.heatmap(cf_matrix_knn_cv/np.sum(cf_matrix_knn_cv), annot=True, fmt='.%'
ax.set_title('K-NN Confusion Matrix\n\n');
ax.set_xlabel('\nPredicted Values')
ax.set_ylabel('Actual Values ');

## Ticket Labels - List must be in alphabetical order
ax.xaxis.set_ticklabels(['0','1'])
ax.yaxis.set_ticklabels(['0','1'])

plt.show()
```

```
NameError                                     Traceback (most recent call last)
Cell In[42], line 6
      3 knn_cv = GridSearchCV(knn,param_grid,cv=5)
      4 knn_cv.fit(x_train_2,y_train)
----> 6 y_pred_knn_cv = knn_cv.predict(X_test_2)
      8 print("Best Score:" + str(knn_cv.best_score_))
      9 print("Best Parameters: " + str(knn_cv.best_params_))

NameError: name 'X_test_2' is not defined
```

In []: